

# Project 1

Chris Harris / Stephan Kostreski

March 12, 2017

## Problem 1 - Warm Up

### Part A

**Generate a sample of 1000 from the normal distribution with a mean of 10 and a standard deviation of 2.25.**

*Hint: Type `?rnorm` into the console to see how to use this function to solve this part of the problem. Also note the arguments required:*

- `n`: number of observations (set to 1000)
- `mean`: the mean value of the distribution (set to 10)
- `sd`: standard deviation: the deviation of values from the mean (set to 2.25)

*NOTE: Don't forget to assign the function result to a variable (using the `<-` operator) so you can refer to it later, like so:*

```
data_dist <- rnorm(n=10, mean=1, sd=2)
```

```
#Generates a list of 1000 values from a normally distributed function with a mean of 10 and a standard deviation of 2.25
data_1000 <- rnorm(n=1000, mean=10, sd=2.25)
View(data_1000)
```

### Part B

**Create a histogram of the data with 10 bins (classes)**


Functions/Arguments needed:

- `hist`
  - `x`: data to be tallied (i.e. the variable where you stored the `rnorm` function call)
  - `xlab`: label for the x-axis
  - `breaks`: number of bins (classes) to categorize data

*Note: R may not create exactly 10 bins based on its algorithm. To gain precision you may provide exact breaks by using `breaks= seq(0,20,by=2)`. This is what we did:*

```
hist(x=data_dist, xlab="Sample Values", breaks=seq(0,20, by=2))
```

```
#Generates a histogram with 10 bins based upon our distribution
hist(x=data_1000, breaks=seq(0,20, by=2))
```



*At this point, we recommend playing with different parameters to obtain colors, different titles and more.  
Hint: use ?hist to learn the parameters and/or google*

## Part C


**Create a normal probability plot and boxplot.**

Functions/Arguments needed:


- boxplot - creates a boxplot
  - x: data to be plotted
- qqnorm - creates normal probability plot
  - y: data to be plotted
- qqline - adds a reference line to the qqnorm plot
  - y: data to add reference line
  - col: line color - col=2 creates a red line, col=3 creates a green line

```
#Create a normal probability plot
qqnorm(y=data_1000, ylab="Values From Sample")
qqline(y=data_1000, col=2)
```

Normal Q-Q Plot



```
#Create a boxplot  
boxplot(data_1000)
```



## Problem 2

### Part A

Create 1000 samples of size 10.

Functions/Arguments needed:

- rnorm - same parameters as before
- replicate
  - n: number of repetitions
  - expr: function to repeat 'n' times
- sample - takes a sample of a particular size from a “population distribution”
  - x: distribution to be sampled from (variable result of rnorm)
  - size: size of sample
  - replace: FALSE (sample without replacement)

```
#Create our overall population to be sampled  
population_dist_norm <- rnorm(1000, mean=10, sd=2.25)
```

```
#Use the sample function to take a random sample with specified sample parameter. Repeat this process u  
samples <- replicate(1000, sample(x=population_dist_norm, size=10, replace=FALSE))
```

### Part B

Calculate the mean for these 1000 and graph the sampling distribution of the sample mean for these samples of size 10.

*There are several ways to accomplish this, but we recommend using a looping construct to take the average of each sample and store each average in a list.*

*Using this method, you should wind up having a list of size 1000. Each element containing the average of that particular sample*

Logical Constructs needed:

- For Loop

Functions/Arguments needed:

- c : creates a list with a specified argument range
- mean
  - x: list (function will take the average of this list)
- hist

As this can be tricky for those new to programming, here is an example of how to accomplish this task:

1. We define our distribution population as so:

```
distribution <- rnorm(n=100, mean=10, sd=2.25)
```

2. We take 10 random samples of size 10 from this distribution

```
samples <- replicate(n=10, sample(x=distribution, size=10, replace=FALSE))
```


3. We create a new list that will eventually hold our sample averages:

```
mean_per_sample <- c(1:10) - Creates a list of size 10
```

4. We then go through each of our individual samples, get the mean of the 10 values in each sample, and store that result in our list.

```
for (index in 1:10) { mean_per_sample[index] <- mean(samples[,index]) }  
#Create a vector (a list) with 1000 elements with default values  
mean_per_sample <- c(1:1000)  
for (index in 1:1000) {  
  mean_per_sample[index] <- mean(samples[,index])  
}  
  
hist(x=mean_per_sample, breaks=10)
```

### Histogram of mean\_per\_sample




### Part C


Create a normal probability plot and boxplot of this sampling distribution.

```
#Same as Problem 1c  
  
qqnorm(y=mean_per_sample)  
qqline(y=mean_per_sample, col=2)
```

Normal Q-Q Plot



```
boxplot(x=mean_per_sample)
```



## Part D: Create the 90, 95 and 99 % confidence intervals.

### 1. Get the mean, standard deviation, and number of samples taken

Functions needed:

- mean
- sd
- length

*#In order to calculate CI, we need the mean and standard deviation of our list of sample means.*

```
a <- mean(mean_per_sample)
```

```
s <- sd(mean_per_sample)
```

*#We also need the number of samples we took*

```
n <- length(mean_per_sample)
```

### 2. Calculate 90% CONFIDENCE INTERVAL:

We are going to use the standard CI formula.

Functions needed:

- qnorm - gives us the z-score given a probability
  - p: probability

Example:

$\text{qnorm}(0.95) = 1.644854$     $\text{qnorm}(0.975) = 1.959964$



```
error <- qnorm(0.95)*(s/sqrt(n)) # 90%
left <- a-error
right <- a+error
```

90% CI: [9.8531994, 9.9246034]

### 3. Calculate 95% CONFIDENCE INTERVAL:

```
error <- qnorm(0.975)*(s/sqrt(n))
left <- a-error
right <- a+error
```

95% CI: [9.8463598, 9.931443]

### 4. Calculate 99% CONFIDENCE INTERVAL:

```
error <- qnorm(0.995)*(s/sqrt(n))
left <- a-error
right <- a+error
```

99% CI: [9.8329923, 9.9448105]

## Problem 3

### Part A

Generate a sample of size 1000 from the exponential with mean 2

Functions/Arguments needed:

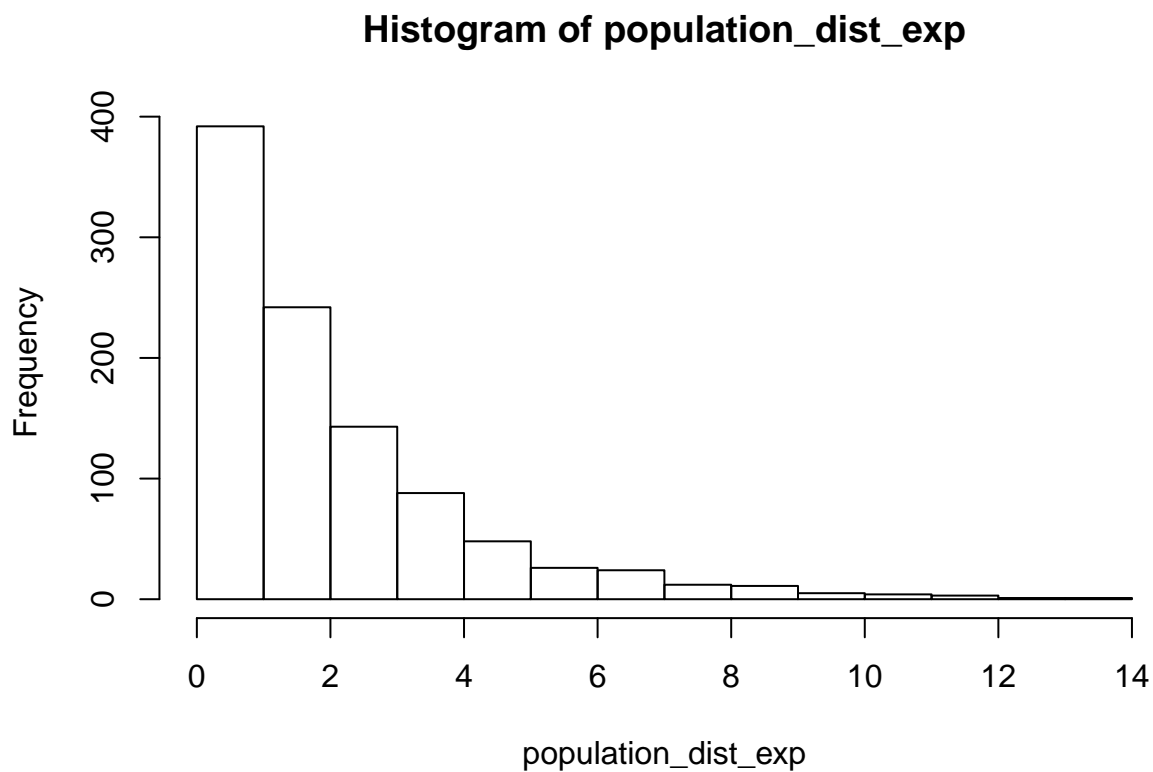
- `rexp` : randomly generates values from the exponential distribution
  - `n`: number of values to be generated
  - `rate`: rate of exponential growth ( $\text{mean} = 1/\text{rate}$ )

```
#Generate population for exponential distribution with mean '2'. Mean = 1/rate. Therefore, rate = 0.5.  
population_dist_exp <- rexp(n=1000, rate=0.5)
```

### Part B

Create a histogram of the data with 10 bins(classes)


```
#Same as Problem 1B  
hist(x=population_dist_exp, breaks=10)
```




### Part C

Create a normal probability plot and boxplot.

```
#Same as Problem 1C  
qqnorm(y=population_dist_exp)  
qqline(y=population_dist_exp, col=2)
```



```
boxplot(x=population_dist_exp)
```



## Part D

1. Create 1000 samples of size 10.

```
samples_size_10 <- replicate(1000, sample(x=population_dist_exp, size=10, replace=FALSE))
```


2. Calculate the mean for these 1000 and graph the sampling distribution of the sample mean for these samples of size 10

```
mean_per_sample_size_10 <- c(1:1000)
for (index in 1:1000) {
  mean_per_sample_size_10[index] = mean(samples_size_10[,index])
}
```

3. Create the histogram of this data with 10 bins.

```
hist(x=mean_per_sample_size_10, breaks=10)
```


**Histogram of mean\_per\_sample\_size\_10**




4. Create a normal probability plot and boxplot of this sampling distribution.

```
qqnorm(y=mean_per_sample_size_10)
qqline(y=mean_per_sample_size_10, col=2)
```

Normal Q-Q Plot



```
boxplot(x=mean_per_sample_size_10)
```



## Part E

1. Create 1000 samples of size 40

```
samples_size_40 <- replicate(1000, sample(x=population_dist_exp, size=40, replace=FALSE))
```


2. Calculate the mean for these 1000 and graph the sampling distribution of the sample mean for these samples of size 40

```
mean_per_sample_size_40 <- c(1:1000)
for (index in 1:1000) {
  mean_per_sample_size_40[index] = mean(samples_size_40[,index])
}
```

3. Create the histogram of this data with 10 bins.

```
hist(x=mean_per_sample_size_40, breaks=10)
```

**Histogram of mean\_per\_sample\_size\_40**




4. Create a normal probability plot and boxplot of this sampling distribution.


```
qqnorm(y=mean_per_sample_size_40)
qqline(y=mean_per_sample_size_40, col=2)
```



Normal Q-Q Plot



```
boxplot(x=mean_per_sample_size_40)
```



5. Create the 90, 95 and 99 % confidence intervals.

```
a <- mean(mean_per_sample_size_40)
s <- sd(mean_per_sample_size_40)
n <- length(mean_per_sample_size_40)

error_90 <- qnorm(0.95)*(s/sqrt(n))
left_90 <- a - error_90
right_90 <- a + error_90

error_95 <- qnorm(0.975)*(s/sqrt(n))
left_95 <- a - error_95
right_95 <- a + error_95

error_99 <- qnorm(0.995)*(s/sqrt(n))
left_99 <- a - error_99
right_99 <- a + error_99
```

90% Confidence Interval [1.99666,2.0303781]

95% Confidence Interval [1.9934303,2.0336079]

99% Confidence Interval [1.9871179,2.0399202]

““